

```
from matplotlib import pyplot as plt
import numpy as np
```

Problem 1 Part A

```
def sigmoid(x):
    return 1.0/(1.0 + np.exp(-x))
```

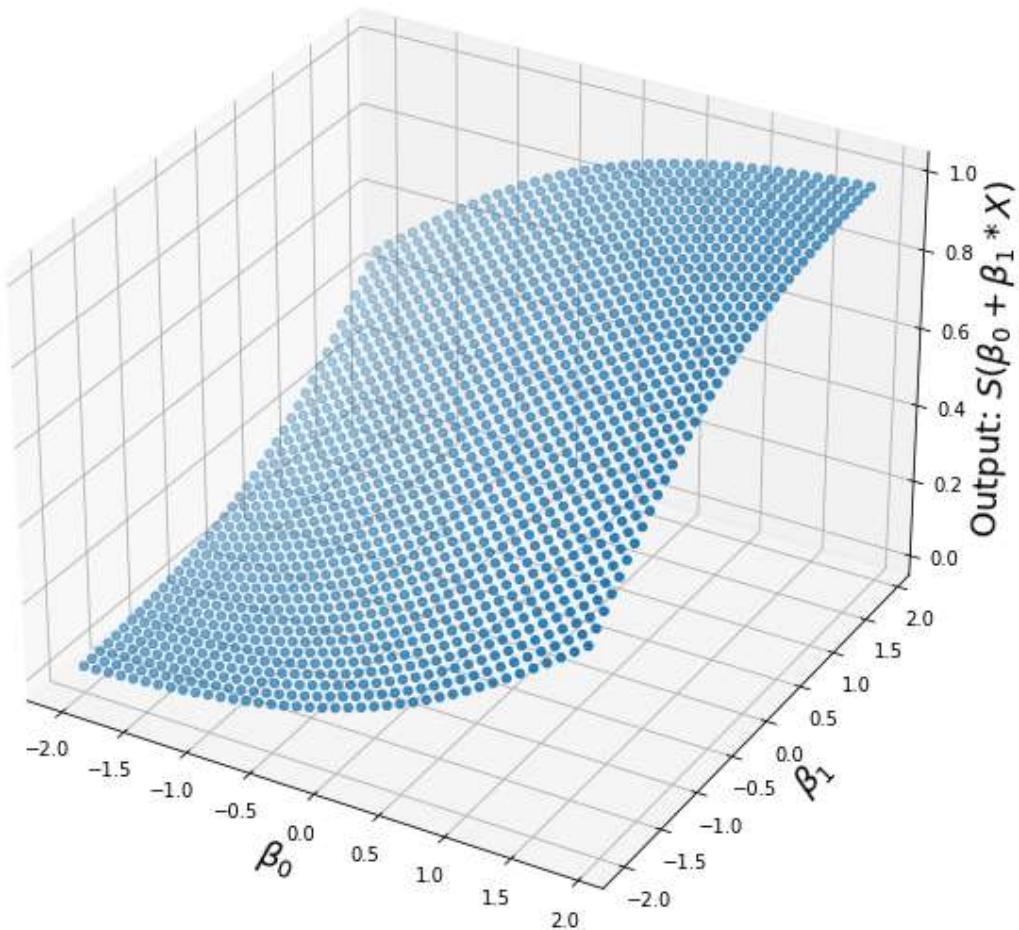
```
beta_0, beta_1 = np.mgrid[-2.0:2.0:0.1, -2.0:2.0:0.1]
```

```
X = 1
sigmoid_input_array = beta_0 + X * beta_1
output_array = sigmoid(sigmoid_input_array)
```

```
plt.rcParams['figure.figsize']=(10, 10)
fig = plt.figure()
ax = fig.add_subplot(projection='3d')
ax.scatter(beta_0, beta_1, output_array)

ax.set_title("3D Plot of Sigmoid Function", fontsize=24)
ax.set_xlabel(r'$\beta_0$', fontsize=18)
ax.set_ylabel(r'$\beta_1$', fontsize=18)
ax.set_zlabel(r'Output: $\beta_0 + \beta_1 * X$', fontsize=18)
plt.savefig("q1_a.jpg")
plt.savefig("q1_a.png")
plt.show()
```

3D Plot of Sigmoid Function



PROBLEM 1B

To prove: $P(Y|X) = \frac{1}{1+e^{-y(\beta_0+\beta_1x)}} \quad \text{if} \quad Y \in \{-1, 1\}$

$$\text{Given: } P(Y=+1|X) = S(\beta_0 + \beta_1 X) = \frac{1}{1+e^{-(\beta_0+\beta_1x)}}$$

$$\begin{aligned} \Rightarrow P(Y=-1|X) &= 1 - S(\beta_0 + \beta_1 X) = \frac{1+e^{-(\beta_0+\beta_1x)} - 1}{1+e^{-(\beta_0+\beta_1x)}} \\ &= \frac{e^{-(\beta_0+\beta_1x)}}{e^{-(\beta_0+\beta_1x)} [e^{(\beta_0+\beta_1x)} + 1]} = \frac{1}{1+e^{(\beta_0+\beta_1x)}} \end{aligned}$$

Thus, to generalize, we use Y as a coefficient in the exponential term

Hence Proved.

To prove: Log likelihood for m data points can be:

$$\ln L(\beta_0, \beta_1) = - \sum_{i=1}^m \ln (1+e^{-y_i(\beta_0+\beta_1x_i)})$$

$$\text{We know from above that } L(\beta_0, \beta_1) = \prod_{i=1}^m \frac{1}{1+e^{-y_i(\beta_0+\beta_1x_i)}}$$

$$\begin{aligned} \Rightarrow \ln(L(\beta_0, \beta_1)) &= \sum_{i=1}^m \ln \left(\frac{1}{1+e^{-y_i(\beta_0+\beta_1x_i)}} \right) = \sum_{i=1}^m (\ln(1) - \ln(1+e^{-y_i(\beta_0+\beta_1x_i)})) \\ &= - \sum_{i=1}^m \ln(1+e^{-y_i(\beta_0+\beta_1x_i)}) \end{aligned}$$

Hence Proved.

Problem 1 Part B

```
def log_likelihood(beta_0, beta_1, X, Y):
    return -np.log(1 + np.exp(-Y * (beta_0 + beta_1 * X)))
```

```
beta_0, beta_1 = np.mgrid[-2.0:2.0:0.1, -2.0:2.0:0.1]
```

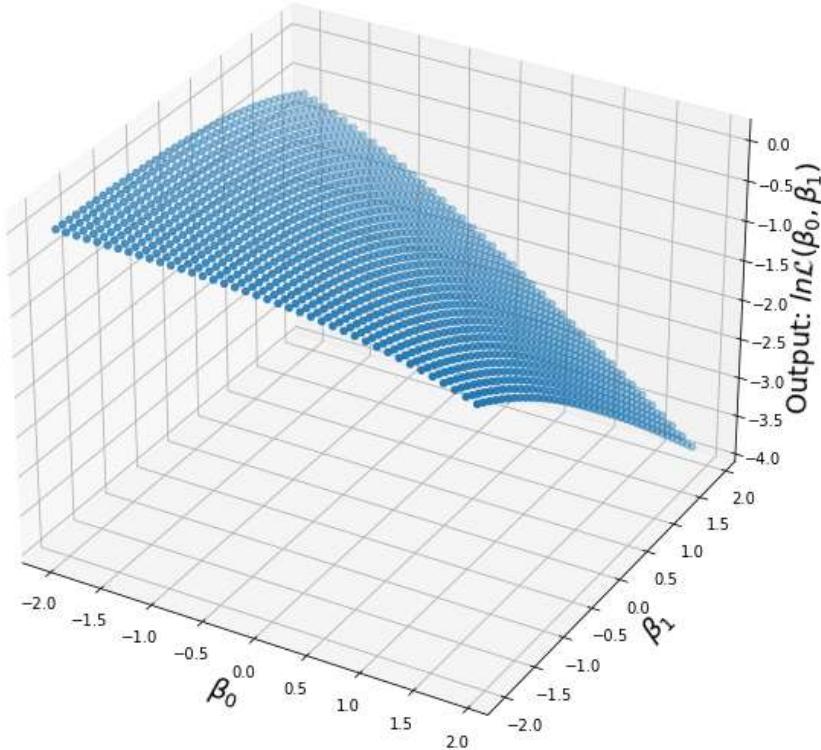
Case 1: $X = 1, Y = -1$

```
X = 1.0
Y = -1.0
output_array = log_likelihood(beta_0, beta_1, X, Y)

plt.rcParams['figure.figsize']=(10, 10)
fig = plt.figure()
ax = fig.add_subplot(projection='3d')
ax.scatter(beta_0, beta_1, output_array)

ax.set_title(r"3D Plot of Log Likelihood Function; $X = 1, Y = -1$", fontsize=24)
ax.set_xlabel(r'$\beta_0$', fontsize=18)
ax.set_ylabel(r'$\beta_1$', fontsize=18)
ax.set_zlabel(r'Output: $\ln \mathcal{L}(\beta_0, \beta_1)$', fontsize=18)
plt.savefig("q1_b1.jpg")
plt.savefig("q1_b1.png")
plt.show()
```

3D Plot of Log Likelihood Function; $X = 1, Y = -1$



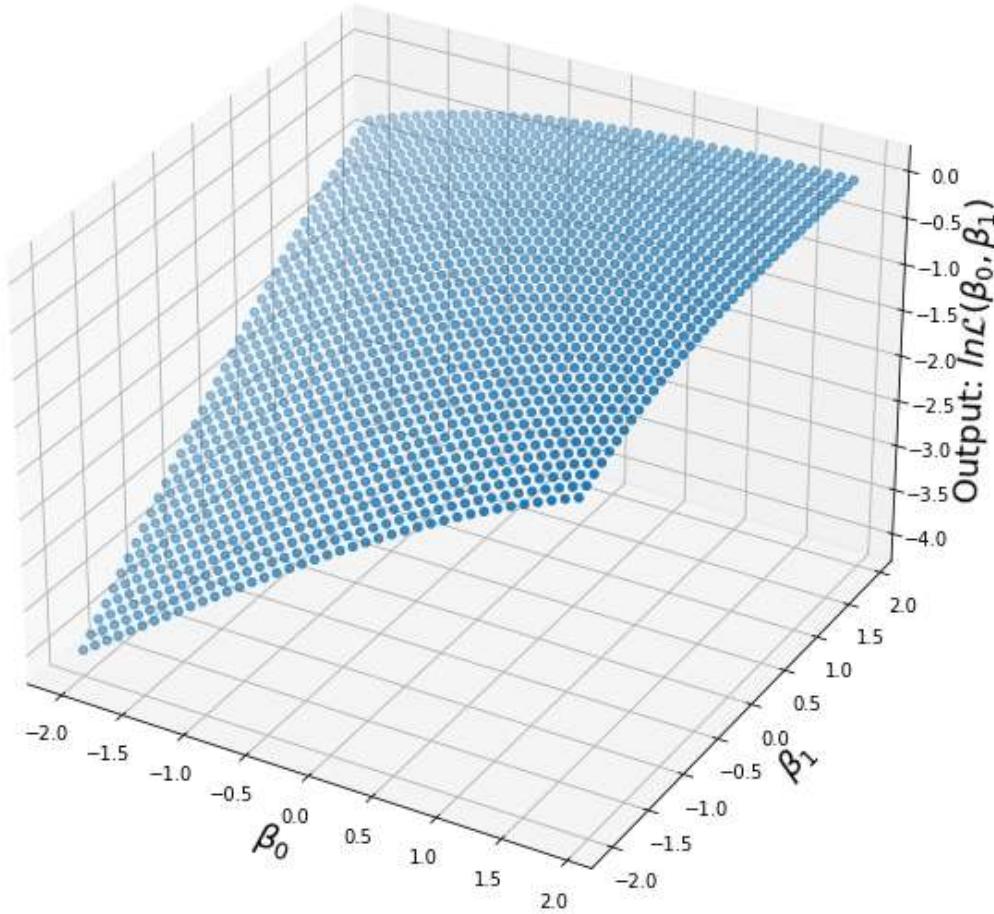
Case 2: $X = 1, Y = 1$

```
X = 1.0
Y = 1.0
output_array = log_likelihood(beta_0, beta_1, X, Y)

plt.rcParams['figure.figsize']=(10, 10)
fig = plt.figure()
ax = fig.add_subplot(projection='3d')
ax.scatter(beta_0, beta_1, output_array)

ax.set_title(r"3D Plot of Log Likelihood Function; $X = 1, Y = 1$", fontsize=24)
ax.set_xlabel(r'$\beta_0$', fontsize=18)
ax.set_ylabel(r'$\beta_1$', fontsize=18)
ax.set_zlabel(r'Output: $\ln \mathcal{L}(\beta_0, \beta_1)$', fontsize=18)
plt.savefig("q1_b1.jpg")
plt.savefig("q1_b1.png")
plt.show()
```

3D Plot of Log Likelihood Function; $X = 1, Y = 1$



Q: Based on the graph, is it possible to maximize this function?

As β_0, β_1 approach their upper, the log likelihood increases. Therefore, the functions are monotonic, and possible to find a global maxima. A caveat is that numerically, the asymptote of this graph is not feasible to reach unless we increase beta 0 and beta 1 infinitely.

PROBLEM 2

Derive the classification rule for threshold 0.5

$$P(Y=1|X) = \frac{1}{1+e^{-y(\beta_0+\beta_1x)}} \geq 0.5 \quad \text{when } y=1$$

$$\Rightarrow 1 \geq 0.5 + 0.5 e^{-(\beta_0+\beta_1x)}$$

$$\Rightarrow 1 \geq e^{-(\beta_0+\beta_1x)}$$

$$\Rightarrow \beta_0 + \beta_1 x > 0 \quad \text{is the classification rule}$$

```

import numpy as np
from matplotlib import pyplot as plt

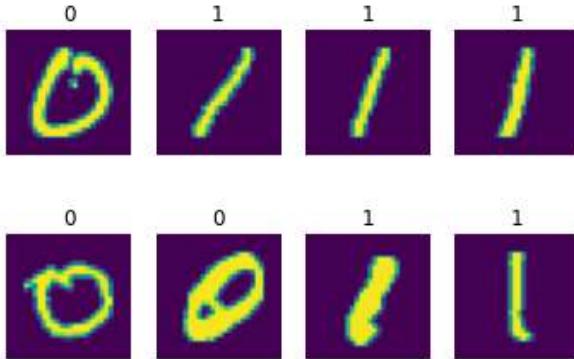
raw_img_data = np.load('data.npy')
label_data = np.load('label.npy')

normalized_img_data = raw_img_data / 255.0
normalized_imshow_array = normalized_img_data.reshape(14780, 28, 28)

index = np.random.randint(0, len(normalized_imshow_array), size=8)
i = 0

fig, axs = plt.subplots(2, 4)
for r in range(2):
    for c in range(4):
        axs[r, c].imshow(normalized_imshow_array[index[i]].reshape(28, 28))
        axs[r, c].set_title(str(label_data[index[i]]))
        axs[r, c].axis('off')
        i+=1

```



```

relabeled_data = np.where(label_data == 1, -1, label_data)
relabeled_data = np.where(relabeled_data == 0, 1, relabeled_data)

```

```
(unique, counts) = np.unique(relabeled_data, return_counts=True)
frequencies_relabeled = np.asarray((unique, counts)).T

(unique, counts) = np.unique(label_data, return_counts=True)
frequencies_original = np.asarray((unique, counts)).T

print("original frequency counts")
print(frequencies_original)
print("after relabeling (1 --> -1, 0 --> 1)")
print(frequencies_relabeled)
```

```
original frequency counts
[[ 0 6903]
 [ 1 7877]]
after relabeling (1 --> -1, 0 --> 1)
[[-1 7877]
 [ 1 6903]]
```

```
train_num_samples = int(0.8 * len(relabeled_data))
test_num_samples = len(relabeled_data) - train_num_samples
print("80-20 Test-Train Split = " + str(train_num_samples) + "-" + str(test_num_samples))

train_indices = np.random.randint(0, len(relabeled_data), size=train_num_samples)
test_indices = np.random.randint(0, len(relabeled_data), size=test_num_samples)
```

```
80-20 Test-Train Split = 11824-2956
```

```
X_train, X_test = normalized_img_data[train_indices], normalized_img_data[test_indices]
Y_train, Y_test = relabeled_data[train_indices], relabeled_data[test_indices]

print("Size of Train vs Test = " + str(len(X_train)) + "-" + str(len(X_test)))
```

```
Size of Train vs Test = 11824-2956
```

```
mu, sigma = 0, 1
d = 28*28
beta_0 = np.random.normal(mu, sigma, 1)[0]
beta_1 = np.random.normal(mu, sigma, d)
```

```
def loss_function(beta_0, beta_1, X, Y):
    """
    :param beta_0, beta_1: Logistic Regression coefficients
    :return: Loss over Training set (X and Y are globals)
    """
    loss_function_sum = 0.0
    m = len(X)
    for i in range(m):
        loss_function_sum += np.log(1 + np.exp(-Y[i] * (beta_0 + np.dot(beta_1, X[i, :]))))
    loss_function_result = loss_function_sum / m
    return loss_function_result
```

```
def compute_gradients(beta_0, beta_1, X, Y):
    """
    :return: Gradient of beta_0, Gradient of beta_1
    """
    d_beta_0_sum = 0.0
    d_beta_1_sum = 0.0
    m = len(X)
    for i in range(m):
        exponent_term = np.exp(-Y[i] * (beta_0 + np.dot(beta_1.T, X[i, :])))
        d_beta_0_sum += Y[i] * exponent_term / (1 + exponent_term)
        d_beta_1_sum += Y[i] * X[i] * exponent_term / (1 + exponent_term)

    d_beta_0 = (-1.0 / m) * d_beta_0_sum
    d_beta_1 = (-1.0 / m) * d_beta_1_sum
    return d_beta_0, d_beta_1
```

```
def evaluate_accuracy(beta_0, beta_1, X, Y):
    correct = 0
    for i in range(len(X)):
        predicted_label_prob = 1.0 / (1.0 + np.exp(Y[i] * (beta_0 + np.dot(beta_1.T, X[i]))))
        prediction = 1 if predicted_label_prob > 0.5 else -1
        label = Y[i]
        is_equal = np.array_equal(prediction, label)
        if is_equal:
            correct += 1

    return (correct/len(X))
```

```

num_iterations = 50
learning_rate = 0.05

training_loss_array, training_accuracy_array = [], []
test_loss_array, test_accuracy_array = [], []
for iter in range(num_iterations):

    loss = loss_function(beta_0, beta_1, X_train, Y_train)
    d_beta_0, d_beta_1 = compute_gradients(beta_0, beta_1, X_train, Y_train)

    beta_0 = beta_0 - learning_rate * d_beta_0
    beta_1 = beta_1 - learning_rate * d_beta_1

    train_accuracy = evaluate_accuracy(beta_0, beta_1, X_train, Y_train)
    test_accuracy = evaluate_accuracy(beta_0, beta_1, X_test, Y_test)
    test_loss = loss_function(beta_0, beta_1, X_test, Y_test)

    training_loss_array.append(loss)
    training_accuracy_array.append(train_accuracy)
    test_loss_array.append(test_loss)
    test_accuracy_array.append(test_accuracy)

if(iter % 5 == 0):
    print("[{:2d}] Accuracy on test set: {:.4f}".format(iter, test_accuracy))

```

```

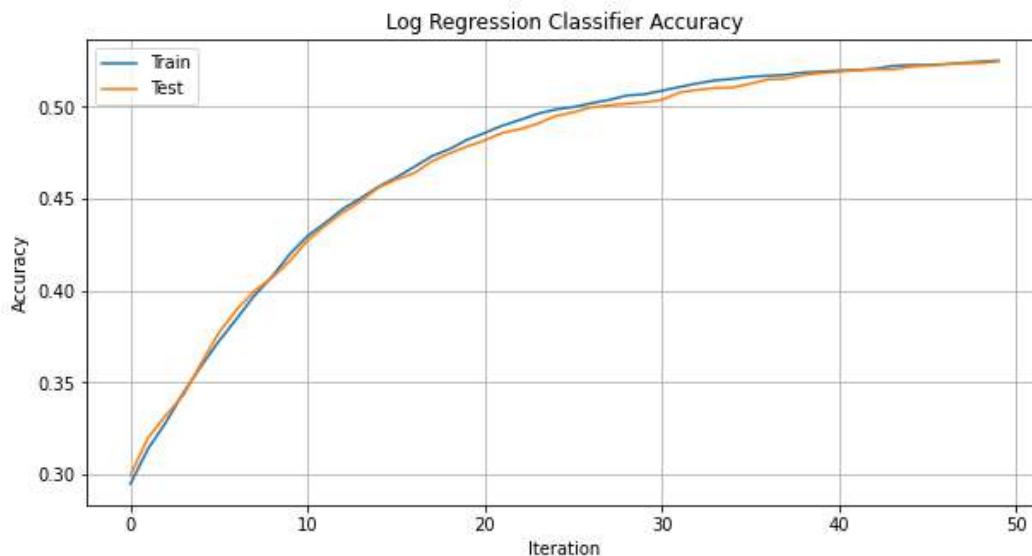
[ 0] Accuracy on test set: 0.2997
[ 5] Accuracy on test set: 0.3772
[10] Accuracy on test set: 0.4273
[15] Accuracy on test set: 0.4604
[20] Accuracy on test set: 0.4817
[25] Accuracy on test set: 0.4970
[30] Accuracy on test set: 0.5037
[35] Accuracy on test set: 0.5125
[40] Accuracy on test set: 0.5189
[45] Accuracy on test set: 0.5223

```

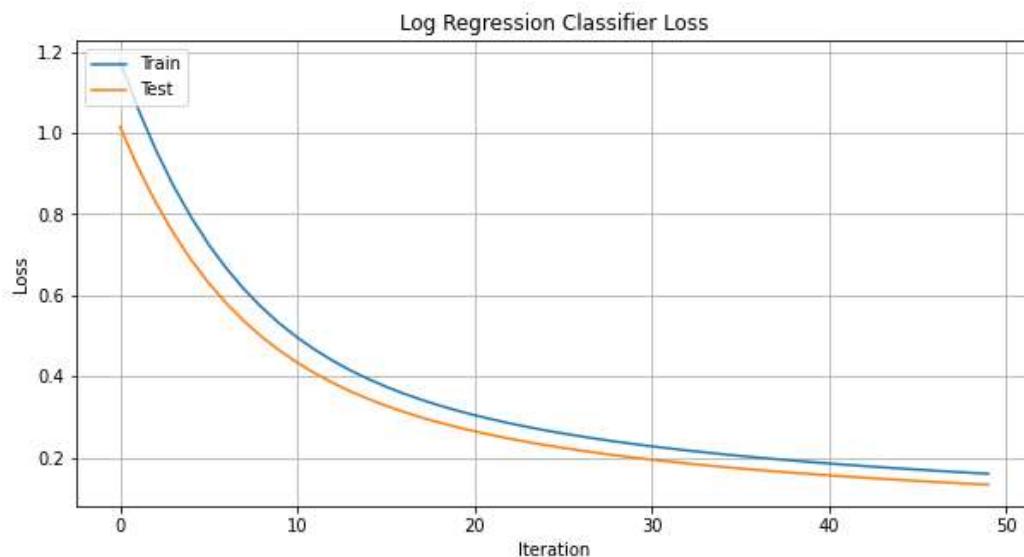
```

plt.rcParams['figure.figsize']=(10, 5)
plt.plot(training_accuracy_array)
plt.plot(test_accuracy_array)
plt.title('Log Regression Classifier Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Iteration')
plt.legend(['Train', 'Test'], loc='upper left')
plt.grid()
plt.show()

```



```
plt.plot(training_loss_array)
plt.plot(test_loss_array)
plt.title('Log Regression Classifier Loss')
plt.ylabel('Loss')
plt.xlabel('Iteration')
plt.legend(['Train', 'Test'], loc='upper left')
plt.grid()
plt.show()
```



PROBLEM 3

Design $P(Y=y|X=x)$ such that i) $P(y=0) = P(y=1) = 0.5$
ii) Classification accuracy of any classifier ≤ 0.9
iii) Accuracy of Bayes optimal classifier is ≥ 0.8

$$P(Y=1|X) = \begin{cases} 0.85 & \text{if } x \leq 0.5 \\ 0.15 & \text{if } x \geq 0.5 \end{cases}$$

$$P(Y=0|X) = \begin{cases} 0.15 & \text{if } x \leq 0.5 \\ 0.85 & \text{if } x \geq 0.5 \end{cases}$$

Are results in $n=1000$ vs $n=100$ any different?

- For Bayes optimal, we get ≈ 0.85 accuracy in both cases, since that is the true probability of the distribution generating a label
- Logistic regression does not get better because data's underlying distribution does not match with what the logistic regression model is trying to fit.

```

import numpy as np
from sklearn.linear_model import LogisticRegression as LogReg
from sklearn.metrics import accuracy_score as accuracy
from sklearn.metrics import roc_curve, roc_auc_score
from matplotlib import pyplot as plt

```

```

def generate_label(x):
    prob_y_equals_one = 0.85 if x < 0.5 else 0.15
    return np.random.choice([1, 0], p=[prob_y_equals_one, 1 - prob_y_equals_one])

def bayes_optimal_classifier(x):
    return 1 if x < .5 else 0

vectorized_generate_labels = np.vectorize(generate_label)
vectorized_bayes_classifier = np.vectorize(bayes_optimal_classifier)

```

```

n_array = [100, 1000]

for n in n_array:
    X_train = np.random.uniform(0, 1, size=n).reshape(-1, 1)
    Y_train = vectorized_generate_labels(X_train).ravel()

    logistic_regression = LogReg()
    _ = logistic_regression.fit(X_train, Y_train)

    X_test = np.random.uniform(0, 1, size=n).reshape(-1, 1)
    Y_test = vectorized_generate_labels(X_test).ravel()

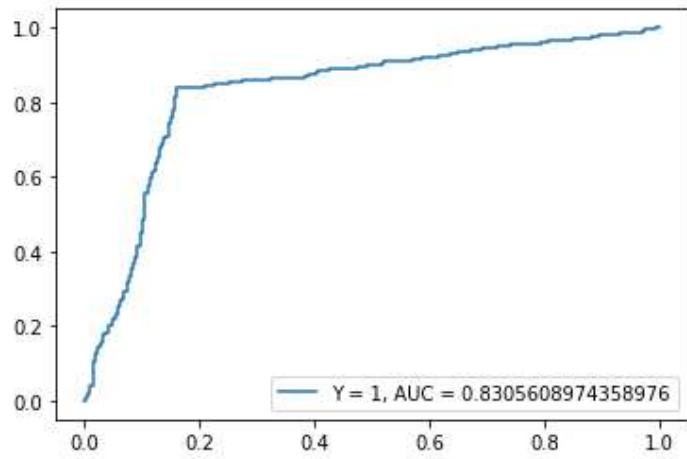
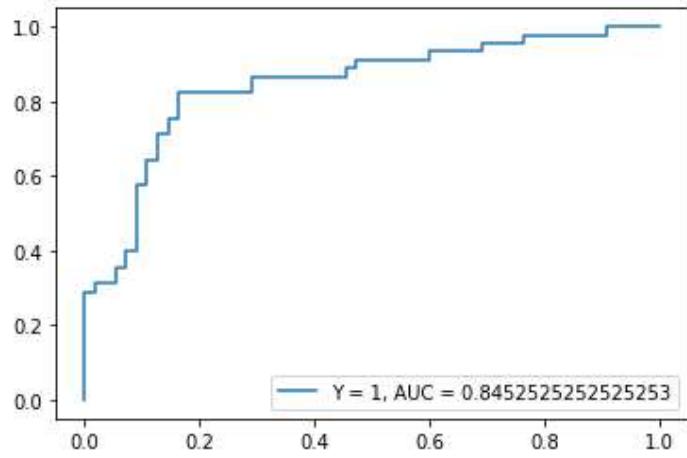
    Y_bayes_test = vectorized_bayes_classifier(X_test)
    Y_pred = logistic_regression.predict(X_test)

    print("===== n = {:5d} =====".format(n))
    print("Logistic Regression Classifier Accuracy : " + str(accuracy(Y_test, Y_pred)))
    print("Bayes Optimal Classifier Accuracy : " + str(accuracy(Y_test, Y_bayes_test)))

    Y_pred_probability = logistic_regression.predict_proba(X_test)[:, 1]
    false_positive_rate, true_positive_rate, _ = roc_curve(Y_test, Y_pred_probability)
    auc = roc_auc_score(Y_test, Y_pred_probability)
    plt.plot(false_positive_rate, true_positive_rate, label="Y = 1, AUC = " + str(auc))
    plt.legend(loc=4)
    plt.show()

===== n = 100 =====
Logistic Regression Classifier Accuracy : 0.7
Bayes Optimal Classifier Accuracy : 0.83
===== n = 1000 =====
Logistic Regression Classifier Accuracy : 0.82
Bayes Optimal Classifier Accuracy : 0.839

```



PROBLEM 4

TODO: Determine an update rule for the centroid c_k of the k -th cluster C_k

Find optimal c_k that minimizes the objective function.

The data x contains p features

4.1

To prove: $\sum_{k=1}^K \sum_{x \in C_k} \sum_{i=1}^p (c_{ki} - x_i)^2$ results in an update rule

where the optimal centroid is the mean of the points in the cluster.

Proof: Within a given centroid (independent & non-overlapping with other centroids), we have:

$$\sum_{x \in C_k} \sum_{i=1}^p (c_{ki} - x_i)^2 \quad \leftarrow \text{update rule}$$

To simplify, first consider the case of a single feature. We can later generalize.

$$\Rightarrow \sum_{x \in C_k} (c_{ki} - x_i)^2 \quad \leftarrow \text{update rule.}$$

$$\text{First derivative set equal to 0} \Rightarrow 2 \sum_{x \in C_k} (c_{ki} - x_i) = 0$$

$$\Rightarrow \sum_{x \in C_k} c_{ki} = \sum_{x \in C_k} x_i$$

$$\text{Assuming centroid } c_k \text{ has } n_k \text{ data points} \Rightarrow n_k c_{ki} = \frac{n_k}{n_k} \sum x_i$$

\Rightarrow

$$c_{ki} = \bar{x} \quad (\text{average of points in centroid})$$

4.2

$$\sum_{k=1}^K \sum_{x \in C_k} \sum_{i=1}^p |c_{ki} - x_i| \quad \text{result in centroid as median of cluster}$$

Following same simplifications as 4.1

$$\Rightarrow \sum_{x \in C_k} \sum_{i=1}^p \text{sign}(c_{ki} - x_i) = 0 \quad (\text{first derivative of } |c_{ki} - x_i|)$$

$$\text{sign}(c_{ki} - x_i) = \begin{cases} -1 & \text{if } c_{ki} < \cancel{x_i} \\ 1 & \text{if } c_{ki} > \cancel{x_i} \end{cases}$$

To minimize ^{set} the LHS function equal to zero,

the # of -1's & +1's must be equal (negating each other & resulting in a 0 sum)

$$\Rightarrow \# \text{ of } x_i \text{'s} > c_{ki} = \# \text{ of } x_i \text{'s} \leq c_{ki}$$

$\Rightarrow c_{ki}$ by definition must be the median of all the points in the ~~centroid~~ cluster.